


Inductive ST-GNN encoder — layer stack

No block carries a parameter whose shape depends on $|V|$ $\rightarrow$ the same ~ 25 k-parameter θ runs on $N = 40, 10$, or 4



GAT layers are vectorized over the $B \cdot H$ per-timestep graphs; TemporalConv mixes only along time, per node.