

Inductive ST-GNN encoder — no parameter shape depends on $|V|$

Size-invariant: the same θ runs verbatim on



$$e_{ij} = \text{LeakyReLU}(a_s^\top W x_i + a_d^\top W x_j) + \log w_{ij}, \quad \alpha_{ij} = \text{softmax}_j(e_{ij}), \quad h_j = \sum_i \alpha_{ij} W x_i, \quad z = [\text{mean}_n h \parallel \text{max}_n h]$$